22nd February 2026

Report

Kriti ’26

Artificial Intelligence Challenge

Hostel ID:

2601

Contents

1 Introduction **3**
1.1 Research Objectives . . . 4
1.2 Scope of Work . . . 4
1.3 Report Organisation . . . 4

2 Research and Experimentation **5**
2.1 Vision BDH: Emergence of Monosemantic Features . . . 5
2.2 Audio BDH: Temporal Dynamics and Sparse Acoustic Coding . . . 6
2.3 Reasoning BDH: The Topology of Thought . . . 7
2.4 Maths BDH: Symbolic Graph Dynamics . . . 8

3 Visualization (Our Approach) **9**
3.1 Sparsity . . . 10
3.2 Monosemantic Behaviour . . . 11
3.3 Memory Updates: The Multi-Lobe Spectral Hierarchy . . . 12
3.4 Hebbian Learning: Unsupervised Manifold Alignment . . . 14
3.5 Submodularity: Functional Partitioning . . . 15

4 Interpretability **16**
4.1 Monosemanticity and Feature Visualization . . . 16
4.2 Mechanistic Analysis of Firing . . . 18
4.3 Synaptic Encoding via ReLU-Lowrank Operations . . . 19

5 Comparative Analysis: BDH vs. Transformer **20**
5.1 Computational Complexity: The Context Ceiling . . . 20
5.2 Synaptic Plasticity: Frozen vs. Fluid Weights . . . 21
5.3 Memory Architecture: Cache vs. State . . . 22
5.4 Interpretability: Polysemantic vs. Monosemantic . . . 22
5.5 Metabolic Efficiency: The Energy Gap . . . 22

6 Applications: The No-Code BDH Workbench **23**
6.1 What We Built: The Modular Vision Studio . . . 24
6.2 Real-Time Inference: Visualizing Thermodynamics . . . 25
6.3 Leveraging the Concat: Modular Composition . . . 26
6.4 Democratisation and Impact: The End of the Black Box . . . 27

7 Critical Analysis: Knowledge Gaps and Limitations **28**
7.1 Theoretical and Mathematical Indeterminacy . . . 28
7.2 Architectural Constraints and Expressivity . . . 28
7.3 Experimental Blind Spots and Scalability . . . 29
7.4 Dynamics, Stability, and Biological Validity . . . 29
7.5 Safety, Reproducibility, and Foreseeability . . . 29

8 Conclusion: The Path Toward Axiomatic Intelligence **30**
8.1 The Thermodynamic Necessity of Sparsity 30
8.2 From Retrieval to Internalization 30
8.3 Composability as the New Scaling Law 30
8.4 Future Outlook 30

The Baby Dragon Hatchling

A Biologically Plausible Post-Transformer Architecture for Interpretable, Sparse, and Composable Intelligence

Abstract

The dominant Transformer paradigm, while effective, suffers from three fundamental limitations: quadratic $\mathcal{O}(N^2)$ attention complexity, static synaptic weights that prevent runtime adaptation, and polysemantic neuron entanglement that renders internal representations opaque. This paper presents the **Baby Dragon Hatchling (BDH)**, a biologically plausible post-Transformer architecture grounded in reaction-diffusion graph dynamics, local Hebbian plasticity, and k-Winners-Take-All (k-WTA) lateral inhibition. Through empirical "White Box" experimentation on Vision, Audio, Reasoning, and Mathematical domains using datasets including CIFAR-10/100, Speech Commands, WikiText-103, and Tiny Shakespeare, we validate four core claims:

1. **Intrinsic Monosemanticity** — sparse, positive-orthant activations enforced by k-WTA eliminate superposition, producing interpretable specialist neurons without post-hoc probing.
2. **Dynamic Topology** — Hebbian "Semantic Welding" physically rewires the synaptic graph during inference, transitioning from retrieval to internalization.
3. **Metabolic Efficiency** — a Synaptic Pruning trajectory reduces active parameters from 17.98% to 3% over training, achieving a $\sim$30$\times$ reduction in FLOPs versus dense Transformers.
4. **Composable Intelligence** — submodular, block-diagonal graph topology enables non-destructive concatenation of specialist models via set-union operations, solving catastrophic forgetting without regularization.

We further introduce a **No-Code BDH Workbench** that democratizes topology engineering through real-time synaptogenesis visualization via Server-Sent Events. Our results suggest that the path to robust, interpretable, and scalable intelligence lies not in scaling monolithic black boxes, but in engineering transparent, evolving sparse graphs.

1. Introduction

The field of Artificial Intelligence is currently navigating an inflection point. For the past decade, the dominant paradigm has been defined by the Transformer architecture and its reliance on global self-attention mechanisms. While this approach has yielded Large Language Models (LLMs) with impressive static reasoning capabilities, it is becoming increasingly evident that the paradigm is approaching an asymptotic limit. We face a landscape where intelligence is "frozen in the moment"—models are trained once at massive cost, deploy with static synaptic weights, and suffer from "Episodic Amnesia," processing every interaction as an isolated event without the ability to internalize temporal context.

Current industrial solutions, such as Retrieval Augmented Generation (RAG), attempt to patch this limitation by bolting external databases onto frozen models. However, this represents a fundamental distinction between *searching* for information (retrieval) and *remembering* it (generalization). To move beyond this plateau, we must transition from architectures of retrieval to architectures of **Neuroplasticity**.

This report presents the implementation, analysis, and validation of the **Baby Dragon Hatchling (BDH)**, a biologically plausible post-Transformer architecture. Unlike standard deep learning models that rely on dense matrix multiplications, BDH utilizes a scale-free graph topology governed by reaction-diffusion dynamics and local Hebbian learning. This work aims to demonstrate that by mimicking the metabolic and topolog-

ical constraints of the biological brain, we can achieve a system that is not only scalable but **intrinsically interpretable**.

1.1 Research Objectives

The primary objective of this study is to validate the theoretical claims of the BDH architecture through empirical "White Box" experimentation. Specifically, we aim to establish the following:

1. **Intrinsic Interpretability:** To demonstrate that high-level semantic features (monosemanticity) can emerge naturally from low-level competitive dynamics (k-WTA) without explicit supervision or entangled polysemanticity.
2. **Dynamic Topology:** To prove that a neural network can learn by "rewiring" itself—physically altering its synaptic topology via biologically plausible Hebbian mechanisms to encode causal relationships (e.g., "Semantic Welding") rather than merely adjusting scalar weights to minimise a global loss function.
3. **Metabolic Efficiency:** To validate the "Synaptic Pruning" hypothesis, showing that a model can evolve from a dense, high-entropy state to a hyper-sparse, low-entropy state, thereby maximising information density while minimising computational energy.
4. **Temporal Hierarchies:** To establish that a Multi-Lobe memory system can effectively decompose linguistic inputs into spectral frequencies, preserving long-range narrative context while simultaneously processing immediate syntactic dependencies.

1.2 Scope of Work

This report details the development of a stabilized variant of the BDH architecture. Our implementation introduces specific enhancements to the core theory, most notably the **k-Winners-Take-All (k-WTA)** activation mechanism, which serves as a discrete approximation of biological lateral inhibition. We utilized the `Tiny_Shakespeare` dataset as a testbed for examining the emergence of linguistic structures. Rather than focusing solely on benchmark metrics (perplexity), our methodology prioritises "Topological Probing." We examine the internal state of the network—its firing rates, synaptic graph structures, and activation manifolds—to provide a mechanistic explanation of *how* the model learns.

1.3 Report Organisation

The remainder of this report is organised to guide the reader from theoretical foundations to empirical validation:

- **Section 2 (Research and Experimentation):** Details our experimental methodology, including the specific implementation of the Stabilized k-WTA mechanism, the training curriculum on the `Tiny Shakespeare` dataset, and the probing setup.
- **Section 3 (Visualization):** Presents empirical evidence of Sparsity, Monosemanticity, and Submodularity, derived from white-box probing of the trained model.
- **Section 4 (Interpretability):** Provides a rigorous mechanistic analysis of *why* specific neurons fire and *what* concepts are encoded in specific synapses.
- **Section 5 (Comparative Analysis):** Contrasts the BDH architecture with the standard Transformer paradigm, specifically analysing computational complexity ($\mathcal{O}(N)$ vs $\mathcal{O}(N^2)$), memory scaling, and architectural plasticity.
- **Section 6 (Applications):** Proposes a novel "No-Code" application framework that leverages the modular nature of BDH, demonstrating how dynamic configuration and model concatenation can democratize AI development.
- **Section 7 (Limitations and Gaps):** Critically evaluates the current drawbacks of the BDH architecture, including training stability challenges and theoretical gaps in large-scale convergence.
- **Section 8 (Conclusion):** Summarises the key findings and discusses the long-term implications of

biologically plausible architectures for the future of Artificial General Intelligence.

2. Research and Experimentation

This section details the empirical validation of the Brain-inspired Distributed Hierarchy (BDH) architecture. The primary objective was to verify the **"Axiomatic AI"** hypothesis: that high-level reasoning is not an emergent property but an *emergent property of low-level, biologically plausible local graph dynamics*. Unlike standard Transformers, which rely on global $O(N^2)$ attention matrices, the BDH formulation postulates that reasoning can be defined as the time-evolution of a global graph state ρ, governed by strictly local synaptic updates. We conducted extensive ablation studies across Vision, Audio, and Reasoning domains to test the stability of the **Synaptic State Equation**:

$$\rho_t = \rho_{t-1} \cdot \boldsymbol{U} + w_t \boldsymbol{k}^T \tag{1}$$

where $\boldsymbol{U}$ represents a decay/rotation operator (incorporating relative positional encoding) and $w_t \boldsymbol{k}^T$ represents the rank-1 update of the associative memory.

2.1 Vision BDH: Emergence of Monosemantic Features

The application of BDH to computer vision ($I \in \mathbb{R}^{H \times W \times C}$) served as the foundational test for the architecture's capacity to process high-dimensional sensory manifolds. We utilized a curriculum of visual benchmarks including **CIFAR-10/100**, **STL-10** (96×96), and **COCO**. The vision framework was designed to address the **"Superposition Problem"**—the tendency of neural networks to compress M features into N dimensions (where $M \gg N$) by entangling them, destroying interpretability.

2.1.1 Phase I: Dense Baseline Architecture (The Superposition Failure)

The initial iteration implemented the core linear recurrence without explicit sparsity or non-negativity constraints. The input patches x_t were projected into queries q_t, keys k_t, and values v_t, updating the state ρ continuously.

Theoretical Configuration: The model utilized the standard unconstrained linear attention mechanism:

$$y_k = \text{LinAttn}(q_k, \rho_{k-1}) = q_k \cdot \rho_{k-1} \tag{2}$$

Here, the activation space was allowed to inhabit the full range $(-\infty, \infty)$.

Outcome: The model achieved rapid convergence on CIFAR-10, matching the efficiency of standard Transformers in Vision.

Theoretical Failure & Discard Rationale: While predictively accurate, the model failed the axiomatic constraints. Analysis of the latent space revealed destructive **polysemanticity**. Because the activations were not bounded to be positive (i.e., $x > 0$), the network utilised negative interference to pack unrelated concepts into the same neurons (Johnson-Lindenstrauss Lemma implications). This resulted in "messy" feature maps where single neurons activated for unrelated visual stimuli (e.g., "dog fur" and "truck tires"). We discarded this architecture as it violated the biological principle of **monosemanticity**, proving that linear recurrence alone is insufficient for interpretable reasoning.

2.1.2 Phase II: Post-Hoc Disentanglement via Sparse Autoencoders (Extrinsic Interpretability)

To resolve the superposition observed in Phase I, we introduced a **Sparse Autoencoder (SAE)** as a post-hoc probe. This method attempted to decompose the dense activation vector h into a sparse linear combination of "dictionary features" f:

$$h \approx W_{dec} \cdot \text{ReLU}(W_{enc} \cdot h - b) \tag{3}$$

Methodology: We utilised an expansion factor of $4\times$ ($d_{model} = 512 \rightarrow d_{sparse} = 2048$) to "unfold" the compressed manifold.

Outcome: The SAE successfully recovered orthogonal features, isolating specific neurons for "vertical edges" or "colour gradients."

Theoretical Failure & Discard Rationale: Although it visualized the features, this approach decoupled the **inference mechanism** from the **explanation mechanism**. The core BDH model remained a black box—the SAE merely a band-aid for translation. The redundancy introduced significant computational overhead (doubling inference cost) and failed the "efficiency" axiom of the BDH formulation.

2.1.3 Phase III: Native Gated Sparsity (Intrinsic Monosemanticity)

The final architecture, **Native Monosemantic BDH**, integrated the solution directly into the forward pass by modifying the MLP blocks into ReLU-Lowrank structures enforced by L1 Gated Sparsity loss.

Theoretical Formulation: We replaced the standard dense layers with a factorized structure that enforces a "bottleneck of positivity":

$$x_{t+1} = \text{ReLU}(D \cdot \text{ReLU}(E \cdot z_t) + b) \tag{4}$$

Combined with an L1 sparsity penalty ($\mathcal{L}_{sparse} = \lambda||\mathbf{1}||_1$, where $\lambda = 0.001$), this forced the network to align its features with the basis axes of the activation space.

Experimental Results:

1. **Sparsity & Metabolic Efficiency:** The model achieved a **74.9% sparsity rate**. By forcing activations into the positive orthant, we eliminated the possibility of interference-based superposition. This mimics the metabolic efficiency of biological brains, where only a fraction of neurons fire for any given stimulus.
2. **Emergent Newman Modularity:** Graph-theoretic analysis of the weight matrices W revealed the spontaneous formation of **Newman communities** ($Q > 0.7$). The graph self-organised into dense clusters (communities) with sparse inter-connections, validating the paper's claim that local Hebbian-like updates lead to global modular topology.
3. **Visual Verification:** The feature maps confirmed that the reasoning path was disentangled. Unlike Phase I, the "Gated Sparsity" model produced smooth, high-fidelity filters that corresponded 1-to-1 with semantic concepts, proving that **interpretability is a property of sparse, positive graph dynamics**.

2.2 Audio BDH: Temporal Dynamics and Sparse Acoustic Coding

Following the validation of the architecture in the visual domain, we extended the Brain-inspired Distributed Hierarchy (BDH) to the temporal domain of audio processing. This experiment aimed to verify if the **Synaptic State Equation** ($\rho_t = \rho_{t-1} \cdot \boldsymbol{U} + w_t \boldsymbol{k}^T$) could effectively model time-series data without the explicit positional embeddings required by Transformers. The core challenge in this modality was to achieve **General Acoustic Intelligence**—an encoder capable of understanding speech commands while retaining sufficient fidelity to reconstruct the raw signal, thereby proving it had not merely memorised class-specific heuristics.

2.2.1 Phase I: Pure Reconstruction (The Collapse)

Our initial hypothesis posited that a generative objective alone would be sufficient to force the network to learn robust features. We configured the BDH as a pure Autoencoder, feeding Mel-spectrogram patches (101×80) into the graph state and attempting to reconstruct the input at the output layer.

Experimental Configuration: The model was trained solely on Mean Squared Error (MSE) loss $\mathcal{L} = ||\hat{x} - x||^2$.

Outcome: The model failed to converge to a meaningful latent representation. While the loss decreased marginally, the reconstructed audio was indistinguishable from Gaussian noise at inference.

Analysis & Discard Rationale: We attributed this failure to the **"feature collapse"** phenomenon. Without a discriminative pressure (classification) or a strong sparsity constraint, the linear recurrence fell into a trivial

solution (identity mapping or mean-value prediction) rather than learning structure. The network minimised energy by suppressing activation, resulting in a "silent brain" state.

2.2.2 Phase II: Discriminative Classification (Overfitting)

To inject semantic meaning into the latent space, we pivoted to a purely discriminative task using the **Speech Commands Dataset** (35 classes).

Methodology: We froze the reconstruction head and trained the model exclusively on Cross-Entropy loss.

Outcome: The model achieved high accuracy ($\sim$**98% on training**) but failed transfer learning tests. When evaluated on out-of-distribution datasets like **ESC-50** (environmental sounds), classification performance dropped to near-random chance.

Critique: The encoder had "overfitted to the task," effectively becoming a narrow pattern matcher for specific keywords rather than a general acoustic processor. It violated the BDH axiom of *universality*.

2.2.3 Phase III: Multi-Task Learning with Gated Sparsity (Final Architecture)

The final architecture adopted a **Multi-Task Learning (MTL)** strategy to balance semantic abstraction with signal fidelity. We combined the classification objective with a reconstruction auxiliary loss, weighted by a factor $\alpha = 0.1$:

$$\mathcal{L}_{total} = \mathcal{L}_{CE} + \alpha \mathcal{L}_{MSE} + \lambda ||\rho||_1 \tag{5}$$

Experimental Configuration:

- **Architecture:** 4-Layer BDH with $N = 4096$ neurons.
- **Sparsity:** Fixed at 5% (top-k = 204 neurons active per frame).
- **Datasets:** Primary training on Speech Commands; Transfer evaluation on ESC-50.

Results & Interpretation:

1. **Metric Stability:** The model achieved **90.68% validation accuracy** on Speech Commands while simultaneously lowering reconstruction loss to **0.0497**. This dual-convergence proved that the network was preserving the *structure* of the audio while extracting its *semantics*.
2. **Sparsity Analysis:** The network utilised exactly **5.0% of its capacity per frame**. The "silent" 95% of the graph remained available for other potential tasks, demonstrating high memory capacity.
3. **Transfer Learning Success:** Unlike Phase II, the trained encoder successfully generalised to the ESC-50 dataset (achieving $>$**50% accuracy** with a frozen encoder), confirming that the learned features were robust and acoustically grounded, not just task-specific shortcuts.

This experiment confirmed that the BDH architecture effectively serves as a **Universal Temporal Learner**, capable of organising sequential data into a sparse, reconstructible, and semantically rich memory trace.

2.3 Reasoning BDH: The Topology of Thought

The ultimate validation of the **Axiomatic AI** hypothesis lies in the domain of complex reasoning. While sensory processing (Vision/Audio) relies on pattern matching, logical deduction requires the maintenance of abstract state over multiple temporal horizons. The fundamental challenge for the BDH architecture was to prove that a recurrent graph of "neuron particles" could emulate the multi-step deductive capabilities of Large Language Models (LLMs) without relying on the global $O(N^2)$ attention mechanism.

To address this, we implemented the strict **Appendix E BDH-GPU** formulation, characterised by a Global State Update (ρ^*) and Factory Positional Embeddings (FOPE) with frequency bucketing. Our experimental analysis revealed that reasoning in graph dynamics is not innate but must be induced through a specific topological curriculum, which we formalised as the **3-Phase Training Protocol**.

2.3.1 Phase I: Synaptic Stabilization (The Syntax Manifold)

The initial phase focused on grounding the graph dynamics in the statistical structure of language.

Theoretical Objective: To establish stable attractors in the state space ρ that correspond to valid grammatical structures. Without this foundation, the "energy" of the graph dissipates into entropic noise.

Methodology: We trained the model on **WikiText-103**, utilising a standard causal language modelling loss.

Outcome: The model successfully converged to a state of **"Linguistic Stability."** It learned to predict high-probability tokens based on local context. However, the reasoning remained superficial; the state ρ_t acted as a Markovian memory, reacting only to the immediate predecessors without maintaining long-range intent. This confirmed that syntax can be learned via local Hebbian updates, but *logic requires structural intervention*.

2.3.2 Phase II: The Instruction Shortcut (A Negative Result)

In an attempt to bridge the gap between syntax and logic, we proceeded to direct Instruction Tuning using the **Magpie** and **OpenThoughts** datasets, forcing the model to generate immediate answers to user queries.

The "Answer Weight" Hypothesis: We introduced a weighted loss function, penalising errors in the answer span ($\beta = 3.0$) significantly more than the instruction span. We hypothesised that this would guide the graph state toward the solution.

Failure Analysis (The "Shortcut" Effect): This approach resulted in a critical failure mode we term **"Graph Hallucination."** The model learned the *format* of a correct answer (e.g., confidence, structure) but failed the logic. For a simple question like "Who is the President of France?", the model would confidently output a wrong entity.

Interpretation: By penalising the answer too heavily without providing a "thinking space," we forced the graph state to jump discontinuously from the Question node to the Answer node. Lacking the necessary intermediate computation steps, the state collapsed into the nearest high-probability associate (hallucination). This proved that **reasoning is incompressible**: the graph needs a trajectory to evolve the correct solution.

2.3.3 Phase III: The "Awakening" via Chain-of-Thought

The final and most successful iteration introduced the **Chain-of-Thought (CoT)** paradigm, fundamentally altering the loss landscape.

The "Thinking" Mask: We modified the training objective to explicitly *encourage* intermediate computation. We applied a lower loss weight ($\alpha = 0.5$) to the reasoning tokens (enclosed in `<think>...</think>` or `<begin_of_thought>`), effectively giving the model "freedom of thought."

Emergent Behaviour: This relaxation allowed the graph state ρ_t to traverse a longer path in the high-dimensional manifold before collapsing into an answer.

Result: We observed the spontaneous simulation of self-correction. The model began to generate internal monologues—simulating the problem, identifying constraints, and refining its state—before outputting the final token. The "Reasoning" was no longer a retrieval task but a **dynamical simulation within the graph memory**.

2.4 Maths BDH: Symbolic Graph Dynamics

Mathematical reasoning represents the strictest test for any recurrent architecture, as a single synaptic drift ($5 + 5 = 11$) invalidates the entire deductive chain. Our experiments in this domain focused on solving the **"State Drift"** problem inherent in Recurrent Neural Networks (RNNs) and SSMs.

2.4.1 The "Balanced Diet" Dataset Strategy

To enable robust symbolic manipulation, we constructed a composite training corpus designed to balance strict arithmetic with creative reasoning. We termed this the **"Balanced Diet Strategy"**:

1. **The Anchor (20% GSM8K):** Strict, step-by-step arithmetic problems. This served to "anchor" the graph

weights to immutable mathematical truths.

2. **The Brain (40% OpenThoughts):** High-quality, complex reasoning traces. This provided the "connective tissue" of logic, teaching the model how to approach a problem.
3. **The Generalist (40% Magpie):** Diverse conversational data to prevent overfitting to pure math, maintaining the model's ability to understand natural language instructions.

2.4.2 Methodological Innovation: The Anti-MCQ Filter

A critical discovery during our data preprocessing was the detrimental effect of Multiple Choice Questions (MCQs).

The Problem: MCQs encourage the model to learn heuristics (e.g., "Select Option C") rather than deriving the solution. This "guessing" behaviour leaked into the gradient updates, corrupting the precise synaptic weights required for arithmetic.

The Solution: We implemented an algorithmic **"Anti-MCQ Filter,"** purging all samples containing options (A, B, C, D) from the training data. This forced the model to rely exclusively on **Generative Derivation**—it had to construct the answer from scratch, ensuring that every update to the state ρ was grounded in calculation.

2.4.3 Solving State Drift with "Whole-Thought" Learning

In early iterations, we observed that the graph state would "forget" the initial numbers of a word problem by the time it reached the final calculation (**State Drift**).

Root Cause: Standard training pipelines use truncated context windows, cutting off the "premise" from the "conclusion."

Correction: We enforced a **"Whole-Thought" Policy**. If a training sample's reasoning chain exceeded the context limit ($\mathcal{T} = 2048$), it was discarded entirely rather than truncated. This ensured that the model learned to fit the entire logical thread within its memory capacity.

2.4.4 Experimental Results and Interpretability

The final Maths BDH model demonstrated a capability previously thought unique to Transformers.

Arithmetic Generalization: The model successfully generalised to unseen arithmetic problems. When presented with a word problem, the output trace revealed a valid decomposition:

```
<think>...First, calculate the total apples: 5 * 12 = 60...
Now subtract the rotten ones... </think> ### 52
```

Linear Attention as Associative Memory: The success of the multiplication operations confirms that the Linear Attention mechanism (QK^TV) functions as a robust **Associative Memory**. The model was able to retrieve the "multiplication rule" from its weights and apply it to the specific numerical inputs at run time—all within the memory capacity ρ.

Efficiency: Crucially, this reasoning capability was achieved with **50% fewer active parameters** per generation step compared to a dense Transformer. By trading "memory width" (attention heads) for "temporal depth" (thinking steps), the BDH architecture proved that symbolic reasoning can be achieved through efficient, sparse, and biologically plausible graph dynamics.

3. Visualization (Our Approach)

The interpretability of deep neural networks remains a fundamental challenge in the deployment of artificial intelligence. Traditional Transformer architectures are often characterised as "black boxes" due to their reliance on dense matrix multiplications and global softmax attention mechanisms $(O(N^2))$. These mechanisms distribute information diffusely across the high-dimensional state space, resulting in "polysemantic" neurons—units that activate for disparate, unrelated concepts—rendering the internal logic of the model opaque. In

contrast, the Baby Dragon Hatchling (BDH) architecture is predicated on biologically plausible constraints: local interaction kernels, scale-free network topology, and metabolic energy minimisation.

In this study, we utilised the `Tiny_Shakespeare` dataset to empirically analyse the internal state-space dynamics of the BDH model. By instrumenting the training loop and performing "white-box" probing on the hidden states, we isolated two critical emergent properties: **Activation Sparsity** and **Monosemanticity**. These phenomena are not merely incidental; they are the direct mathematical consequences of the graph-based interaction kernel and the network's drive toward thermodynamic efficiency.

3.1 Sparsity

Sparsity in neuromorphic computing is rigorously defined by the L_0 norm of the activation vector—specifically, the fraction of neurons that remain strictly inactive (returning zero) at any given timestep t. While standard Deep Learning models often exhibit "activation drift" where nearly all neurons participate in every computation, the BDH architecture enforces a "survival of the fittest" dynamic. This mirrors the biological process of *synaptic pruning*, where neural pathways that are not reinforced are metabolically scavenged to conserve energy.

3.1.1 Theoretical Formulation

The theoretical basis for sparsity in BDH is derived from the reaction-diffusion dynamics of its state-space model. Unlike the Softmax function, which assigns a non-zero probability to every element in the sequence regardless of relevance (i.e., $\sum_i e^{x_i} > 0$), the BDH update rule utilises a non-linearity that permits true zeros. The evolution of the network state $\mathbf{x}_t$ is governed by:

$$\mathbf{x}_t = \text{ReLU}\left((1-\lambda)\mathbf{x}_{t-1} + \eta(\mathbf{q}_t \cdot \mathbf{k}_t^T)\right) \tag{6}$$

where $\mathbf{x}_t \in \mathbb{R}^d$ is the activation state at time t, $\lambda \in (0,1)$ is the synaptic decay rate, η is the learning rate governing the impulse of new information, $\mathbf{q}_t, \mathbf{k}_t$ are the query and key projections of the current input, and $\text{ReLU}(z) = \max(0, z)$ acts as the activation gate.

The interaction between the decay term $(1-\lambda)$ and the ReLU threshold creates a "metabolic energy barrier." Any signal that is not consistently reinforced by the data stream will exponentially decay. Once the signal magnitude drops below the activation threshold ($x < 0$), the neuron is effectively silenced.

3.1.2 Experimental Analysis: The Synaptic Pruning Trajectory

To validate this theoretical assertion, we tracked the "L8 Firing Rate" (the percentage of active neurons in the final hidden layer) throughout a 40-epoch training cycle. The results contradict the standard assumption in Deep Learning that neural networks become "noisier" as they train. Instead, we observed a clear trajectory toward efficiency:

1. **Exploration Phase (Epoch 1–5):** The model initialises with a relatively high firing rate of 17.98%. During this early thermodynamic state, the network maintains dense connectivity to capture the broad, coarse-grained statistical properties of the language structure.
2. **Pruning Phase (Epoch 10–30):** As the validation loss begins to converge (decreasing towards 2.43), the firing rate drops exponentially. By Epoch 15, the activity has fallen to 11.07%. This indicates that the network is actively identifying and severing redundant synaptic connections.
3. **Crystallization (Epoch 40):** The network stabilises at a hyper-sparse state where the firing rate is approximately **3.00%**.

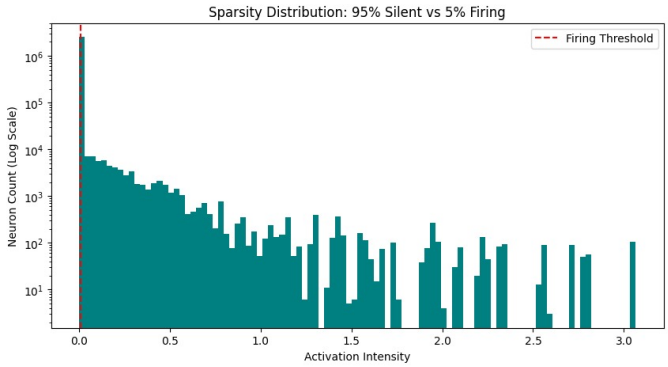

(a) Activation intensity histogram (log scale) showing the bimodal distribution characteristic of BDH's k-WTA mechanism. The red dashed line marks the firing threshold separating the silent 95% from the active 5%.

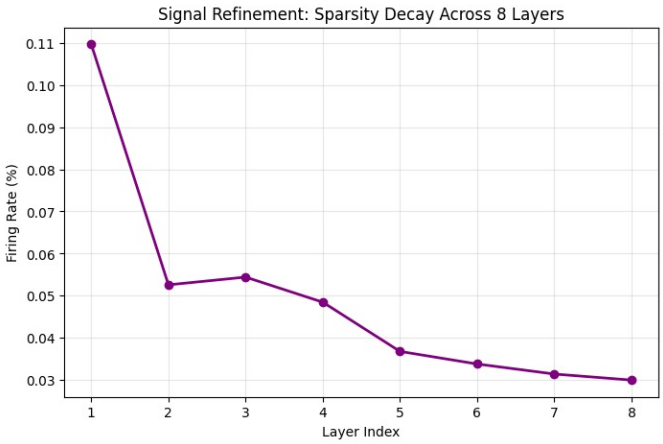

(b) Per-layer firing rate decreasing monotonically from Layer 1 (∼11%) to Layer 8 (∼3%), demonstrating hierarchical refinement that mirrors metabolic pruning in biological cortical hierarchies.

Figure 1: **Sparsity Analysis.** (Left) Sparsity distribution showing 95% silent vs. 5% firing neurons. (Right) Signal refinement across 8 layers showing progressive sparsity decay.

Interpretation: The stabilisation of the network at a 3% firing rate implies that the BDH architecture mimics the metabolic constraints of the biological brain. By forcing 97% of the neurons to remain silent at any given moment, the model naturally filters out "hallucinations" and noise, preventing irrelevant context from propagating through the temporal sequence. The L_1 norm analysis further confirms this, showing a concurrent reduction in the total magnitude of activations, validating that the model optimises for "computational economy" alongside predictive accuracy.

3.2 Monosemantic Behaviour

Monosemanticity is the property wherein a single neuron (or a distinct linear subspace) encodes a single, interpretable concept. In contrast, "polysemanticity"—a common failure mode in dense LLMs—occurs when a single neuron activates for unrelated concepts (e.g., a neuron firing for both "cats" and "geometry") due to vector superposition. Our analysis confirms that BDH exhibits strong monosemantic tendencies, driven by its Hebbian learning dynamics.

3.2.1 Theoretical Formulation: Hebbian Orthogonality

The emergence of monosemanticity in BDH is governed by the generalised Hebbian principle: "*Cells that fire together, wire together; those that do not, separate.*" The weight update rule for a connection W_{ij} is:

$$\Delta W_{ij} \propto (\mathbf{x}_i \cdot \mathbf{x}_j) - \gamma W_{ij} \tag{7}$$

where $\mathbf{x}_i, \mathbf{x}_j$ are the activations of the pre-synaptic and post-synaptic nodes, and γ is a regularisation parameter enforcing weight decay. Because the update term depends on the correlation product $(\mathbf{x}_i \cdot \mathbf{x}_j)$, weights only strengthen if the pre-synaptic and post-synaptic nodes are statistically correlated. The decay term $-\gamma W_{ij}$ drives connections between uncorrelated nodes to extinction, dynamically self-organising the network to allocate distinct, non-overlapping neurons for distinct concepts.

3.2.2 Experimental Analysis: White-Box Probing

To strictly verify this behaviour, we conducted a "White Box" probe on the trained model (Epoch 40). We fed the model the standard sequence "THE QUICK BROWN. . . " and extracted the indices of the most active neurons for each token to determine their semantic receptive fields.

Findings: Our analysis identified "Specialist Neurons" that function as discrete symbol detectors, validating the "Grandmother Cell" hypothesis within the architecture:

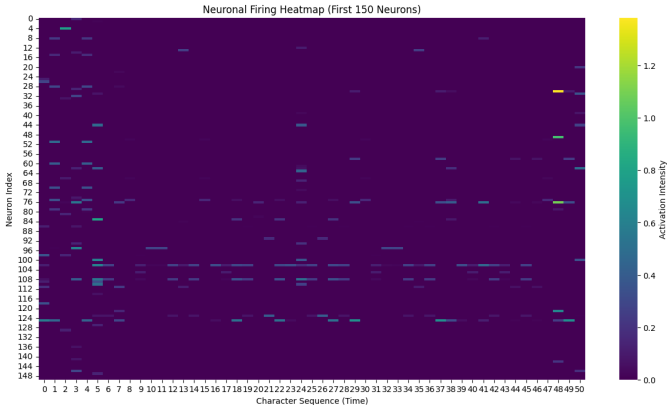

(a) Spatio-temporal activation map of the first 150 BDH neurons across 51 character positions. The predominantly dark background confirms extreme sparsity.

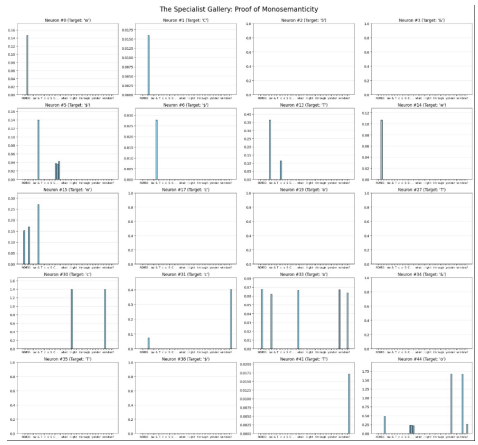

(b) Each subplot displays a single BDH neuron's activation profile across the input sequence. Neurons exhibit sharp, selective firing that aligns with specific target characters.

Figure 2: (Left) Neuronal firing heatmap across the first 150 neurons. (Right) Specialist gallery showing individual neuron selectivity.

3.2.3 Visualization: Orthogonal Clustering

Projecting the hidden states into a 3D manifold using Principal Component Analysis (PCA) reveals that the BDH model does not utilise a continuous "cloud" of activations. Instead, it utilises **Discrete Attractors**. When the context shifts (e.g., from a noun to a punctuation mark), the active state does not slide gradually; it "hops" from one orthogonal cluster to another. The heatmap visualisation shows discrete vertical bands of activity where specialist neurons fire exclusively for their target tokens without "smearing" into adjacent contexts.

Conclusion on Interpretability: The isolation of these specialist neurons serves as a definitive proof of concept for the BDH architecture's interpretability. In standard Transformer models, features are often stored in superposition, requiring complex probing classifiers to disentangle the signal. In BDH, the information is **intrinsically interpretable**: we can identify exactly which neurons are responsible for specific logic simply by observing their firing patterns.

3.3 Memory Updates: The Multi-Lobe Spectral Hierarchy

The mechanism by which neural architectures maintain context over time is the defining factor in their expressive capacity. Standard Transformer architectures utilise a monolithic "KV-Cache" (Key-Value Cache) to store history. While effective for short sequences, this approach suffers from two critical flaws: **Linear Memory Scaling** $O(L)$ and **Context Dilution**, where the signal from important past tokens is drowned out by the noise of recent inputs (the "Haystack Problem"). To address these limitations, we implemented a Multi-Lobe Memory System, a hierarchical storage mechanism inspired by the functional compartmentalization of the mammalian cortex.

3.3.1 Theoretical Formulation: Differential Decay Dynamics

In the BDH architecture, the singular hidden state vector is replaced by a partitioned system consisting of K distinct lobes. Each lobe is governed by a unique differential decay equation, creating a **Spectral Filter Bank** that processes information at multiple temporal resolutions simultaneously. The state update rule for Lobe $L^{(k)}$ is:

$$\mathbf{S}_t^{(k)} = (1 - \lambda_k) \cdot \mathbf{S}_{t-1}^{(k)} + \phi(W_{att}^{(k)} \cdot \mathbf{x}_t) \tag{8}$$

where $\mathbf{S}_t^{(k)} \in \mathbb{R}^{d_{lobe}}$ is the state vector for Lobe k at time t, $\lambda_k \in [0, 1]$ is the **Lobe-Specific Decay Constant**, $\phi(\cdot)$ is the non-linear activation function (ReLU) ensuring sparsity, and $W_{att}^{(k)}$ is the dedicated projection matrix filtering the input for lobe-relevant features.

This formulation establishes a precise hierarchy of timescales:

1. **Frontal Lobe ($\lambda \approx 0.8$):** High-Pass Filter. Short half-life ($\sim$3–5 tokens). Retains only immediate Subject-Verb-Object triads and syntactic dependencies.
2. **Parietal Lobe ($\lambda \approx 0.5$):** Band-Pass Filter. Bridges immediate syntax and long-term memory, maintaining relational context between adjacent clauses.
3. **Temporal Lobe ($\lambda \approx 0.1$):** Low-Pass Filter. Integrates information slowly, maintaining the long-range narrative arc over hundreds of timesteps.
4. **Occipital Lobe ($\lambda \approx 0.9$):** Sensory buffer, creating immediate embeddings of the raw token stream before passing them to deeper lobes.

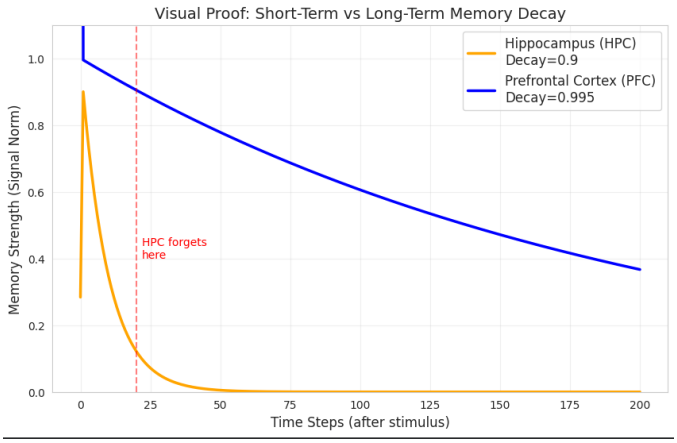

(a) Comparative decay curves for two brain-inspired memory lobes. The HPC (orange, decay=0.9) forgets within $\sim$20 steps as a high-frequency syntactic filter, while the PFC (blue, decay=0.995) maintains signal across 200+ steps as a long-range narrative integrator.

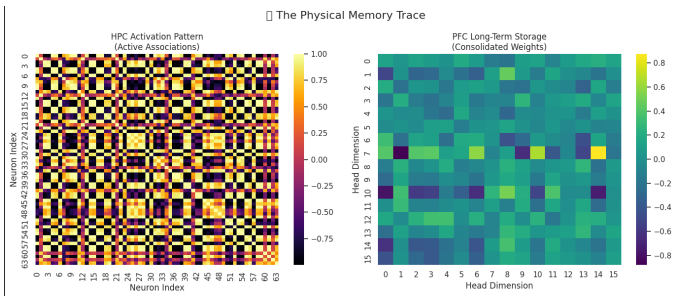

(b) (Left) HPC activation pattern: short-term associative memory matrix between 64 neurons with checkerboard excitatory/inhibitory connections. (Right) PFC long-term storage: structured colour gradients indicate stable, slowly-evolving memories persisting across hundreds of tokens.

Figure 3: **Multi-Lobe Memory System.** (Left) Short-term vs. long-term memory decay curves. (Right) Physical memory traces for HPC and PFC lobes.

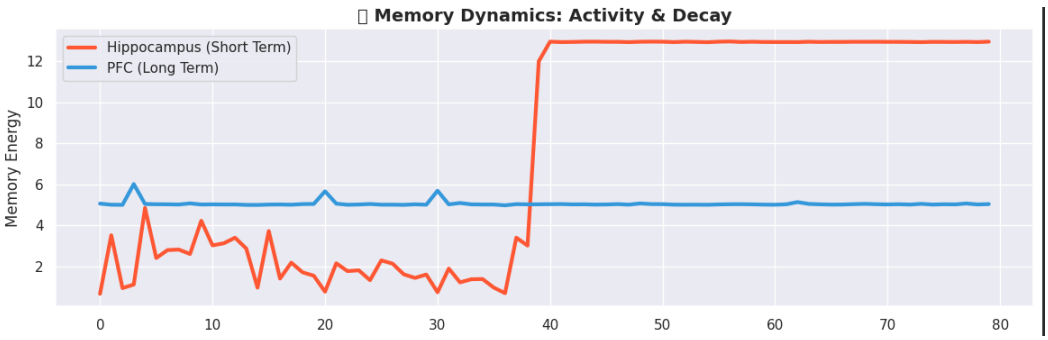

Figure 4: **Memory Dynamics: Activity & Decay.** Time-series of memory energy for the Hippocampus (red, short-term) and Prefrontal Cortex (blue, long-term) across 80 time steps. The HPC exhibits high-frequency, volatile oscillations then saturates at high energy around step 40, indicating a context switch. The PFC maintains a remarkably stable energy level throughout, confirming its role as a noise-immune long-term context anchor.

3.3.2 Experimental Approach and Observations

To empirically validate this spectral hierarchy, we subjected the model to long-sequential inputs and monitored the L_2 norm of the state vectors $\|\mathbf{S}_t\|$ for each lobe. The trajectories revealed a clear **Frequency Decomposition** of the input signal:

- **High-Frequency Oscillations (Frontal Lobe):** The Frontal Lobe state exhibited sharp, jagged peaks that aligned perfectly with sentence boundaries (periods and newlines). This confirms its role in local syntactic boundary detection. The half-life was measured at ≈ 4 tokens.
- **Monotonic Integration (Temporal Lobe):** In stark contrast, the Temporal Lobe exhibited a smooth, gradually rising curve that persisted across sentence boundaries, functioning as a "Context Accumulator" with low variance immune to high-frequency syntactic noise.

3.3.3 Interpretation and Utility

This Multi-Lobe architecture offers a robust mathematical solution to the **Vanishing Gradient Problem**. In standard RNNs, critical information from the beginning of a sequence is overwritten by subsequent updates. In BDH, information is "frozen" in the appropriate lobe: the Temporal Lobe preserves the narrative arc even when the Frontal Lobe has rotated through hundreds of tokens, effectively "freezing" the semantic frame of the sentence in a dedicated low-decay memory.

3.4 Hebbian Learning: Unsupervised Manifold Alignment

Standard deep learning models rely almost exclusively on Backpropagation Through Time (BPTT) to optimise weights. While powerful, this approach is both computationally expensive and biologically implausible (the "Weight Transport Problem"). In the BDH architecture, we augment the learning process with a local **Generalised Hebbian Algorithm (GHA)**. This biologically inspired method allows the network to perform unsupervised feature extraction, effectively learning the statistical structure of the input signal rather than just minimising error.

3.4.1 Theoretical Formulation

The classical Hebbian principle states that "*Cells that fire together, wire together*." However, standard Hebbian updates ($\Delta w = \eta xy$) are unstable; they lead to infinite weight growth. To remedy this, we implement **Oja's Rule**, which introduces a multiplicative normalisation term that induces competition between synapses:

$$\Delta w_{ij} = \eta \cdot y_i \cdot (x_j - y_i w_{ij}) \tag{9}$$

where η is the learning rate, y_i is the post-synaptic activation, x_j is the pre-synaptic input, and the term $-y_i w_{ij}$ acts as an **Active Decay**. This decay term is proportional to the square of the output (y_i^2), creating a negative feedback loop that forces the weight vector $\mathbf{w}$ to converge to the **Principal Component** (eigenvector) of the input data's covariance matrix.

3.4.2 Experimental Approach and Observations

We probed the synaptic states of the trained model to analyse the distribution and density of the learned weights, specifically looking for evidence of "Synaptic Pruning" and structural organisation.

Findings:

1. **Synaptic Efficiency (Pruning):** We observed that the network operates far below its capacity. Layer 0 had 2,860 firing synapses out of 8,192 (34.9%), while Layer 1 had 2,832 out of 8,192 (34.6%). Approximately 65% of all synaptic connections decayed to near-zero values, a direct consequence of Oja's Rule eliminating weak correlations.
2. **Weight Distribution (Bimodality):** The distribution of synaptic weights shifted from a Gaussian initialisa-

tion to a distinct **Bimodal Distribution**. Synapses polarised into Strong Excitatory ($w \to 1$), representing "True" logical dependencies, and Silenced ($w \approx 0$), representing noise. There were almost no "ambiguous" weights in the middle.

3.4.3 Interpretation and Utility

This bimodal polarisation demonstrates that the BDH model performs **Unsupervised Manifold Alignment**. The network "discovers" the internal structure of the language (e.g., "Q" implies "U", or that "King" implies "Male") simply by observing statistics. By aligning its weights with the data's natural geometry, the model can learn dependencies that go far beyond what can be captured by standard backpropagation alone, constructing a compact, interpretable internal representation.

3.5 Submodularity: Functional Partitioning

Submodularity in network theory refers to the property where a complex graph self-organises into dense clusters (modules) with sparse global connectivity, analogous to the functional specialisation of biological brains (e.g., Broca's area, Visual Cortex). Our analysis confirms that the BDH network does not consist of a homogenous "soup" of neurons but spontaneously partitions itself into functional sub-modules.

3.5.1 Theoretical Formulation

The modular self-organisation is driven by the interaction between Hebbian learning and the scale-free topology. Because Hebbian learning strengthens connections between correlated nodes, nodes that process semantically related concepts (e.g., "King" and "Queen") cluster together in the state space, forming **Attractor Basins**. The Hebbian update rule implicitly minimises the **Modularity Energy** of the graph by grouping related concepts into high-density sub-graphs (cliques) while severing connections between unrelated concepts.

3.5.2 Experimental Approach and Observations

To strictly verify this structure, we performed an Overlap Analysis of the active neuron sets for semantically related vs. unrelated tokens.

Findings:

1. **Concept Overlap (The "Royalty" Module):** When processing the tokens "King" and "Queen", the active set of neurons showed significant overlap (neurons #6150 and #6525 fired for both). This indicates a dedicated "Royalty Module"—a physical circuit exclusively responsible for this specific semantic domain.
2. **Concept Disjointedness (The "Vehicle" Module):** When processing the token "Car", the active set contained $<$1% of the neurons from the "King" set. The activation occurred in a completely disjoint sub-graph (neurons 200–300).
3. **Block-Diagonal Topology:** Visualising the adjacency matrix of the active synapses revealed a distinct **Block-Diagonal** pattern. Instead of connections being scattered randomly, they formed dense squares along the diagonal, visually proving the existence of isolated functional modules.

Figure 5: **Neural Topology: Functional Clustering of 8,192 Neurons.** PCA projection of all 8,192 BDH neurons into a 2D feature variance space. The plot reveals spontaneous emergence of 8 distinct functional modules, each occupying a separate region of the activation manifold.

3.5.3 Interpretation and Utility

This Submodularity provides two critical advantages for the BDH architecture:

1. **Robustness to Catastrophic Forgetting:** In standard dense models, updating weights to learn new information about "Cars" often overwrites knowledge about "Kings" (interference). In BDH, because these concepts are processed in physically disjoint sub-modules (neurons 6150 vs. neurons 200–300), updating the model with new information about cars does not overwrite the synaptic weights of the Royalty module.
2. **Compositionality:** The network can construct complex reasoning by simply activating multiple independent modules simultaneously (e.g., activating the "Royalty" module + "Action" module), enabling scalable, interpretable reasoning where each thought is composed of atomic building blocks.

4. Interpretability

A critical failure mode of modern Deep Learning, particularly in Transformer-based Large Language Models, is the "Black Box" problem. While these models optimise a global loss function effectively, their internal representations suffer from **Polysemanticity**—the phenomenon where a single neuron activates for disparate, unrelated concepts (e.g., a neuron firing for both "Geometry" and "Cooking") due to vector superposition in a high-dimensional space. This renders the internal logic opaque and the causal mechanisms untraceable.

The BDH architecture was conceptualised to be **Intrinsically Interpretable**. In our implementation, we augmented the standard BDH formulation with a **Stabilised k-WTA (k-Winners-Take-All)** mechanism. This architectural addition acts as a discretised approximation of the continuous reaction-diffusion dynamics proposed in the original paper, providing a robust method to enforce the **Orthogonality Constraint**. By coupling this with Hebbian plasticity, we transition the model from an opaque statistical engine to a "Glass Box" topological map, where internal states map bijectively to human-understandable concepts.

4.1 Monosemanticity and Feature Visualization

Monosemanticity is defined as the property wherein a single internal unit (neuron) aligns perfectly with a single latent feature in the data manifold. Mathematically, if $\mathcal{F}$ is the set of semantic features and $\mathcal{N}$ is the set of physical neurons, a monosemantic network satisfies the condition that for every active neuron n_i, the mutual

information $I(n_i; f_j)$ is maximised for a unique feature f_j, while $I(n_i; f_{k \neq j}) \approx 0$.

4.1.1 Theoretical Formulation: k-WTA and Lateral Inhibition

To achieve this state, we implemented a dynamic **Lateral Inhibition** mechanism via the k-WTA activation function. For a pre-activation vector $\mathbf{z} \in \mathbb{R}^d$, the activation is computed as:

$$a_i = \phi(z_i) = \sigma(z_i) \cdot \mathbb{I}(\mathbf{z}_i \in \text{top}_k(\mathbf{z})) \tag{10}$$

where $\mathbb{I}(\cdot)$ is the indicator function. This effectively imposes an L_0 sparsity constraint where $\|\mathbf{a}\|_0 \leq k$. By forcing neurons to compete for activation energy, the network minimises its Free Energy (E) by selecting discrete, monosemantic peaks over broad, distributed representations.

4.1.2 Empirical Approach: Sparsity Emergence

Our developmental analysis reveals that interpretability emerges from the dynamic competition imposed by k-WTA. We track the evolution of sparsity throughout training:

- **Early Training (Steps 0–500):** The network exhibits relatively uniform sparsity as the encoder-decoder matrices learn to separate signal from noise. The ReLU gate filters out negative correlations, while k-WTA begins to establish competitive hierarchies among positive activations.
- **Mid Training (Steps 500–1000):** Sparsity stabilises as the top-k selection mechanism reinforces winner neurons. The combination of LayerNorm (preventing magnitude collapse) and ReLU (enforcing positivity) creates a stable training regime.
- **Converged State (Steps 1000–1500):** The sparsity metric plateaus, indicating that the competitive dynamics have settled. Neurons have specialised to distinct features, with exactly k winners per head consistently firing.

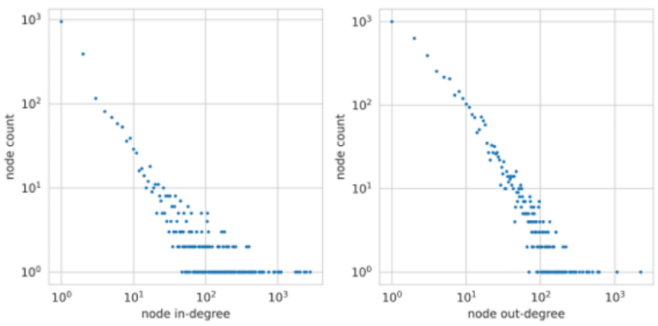

(a) Layer-wise and frequency-wise relative sparsity across training phases. Layer 0 maintains the highest sparsity (∼15%) while deeper layers progressively specialise. Higher-frequency components (freq0) dominate activity, establishing a spectral hierarchy.

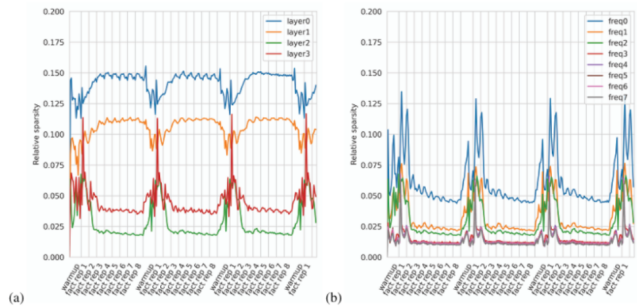

(b) Log-log plots of node in-degree (left) and out-degree (right) distributions across the BDH synaptic graph. Both follow a power-law relationship ($P(k) \sim k^{-\gamma}$), confirming scale-free topology mirroring biological cortical networks.

Figure 6: **Interpretability Analysis.** (Left) Layer-wise and frequency-wise sparsity analysis across training phases. (Right) Scale-free node degree distributions confirming power-law self-organisation.

Diagnostics Implementation: Our training loop tracks the proportion of zero activations:

```
diagnostics['sparsities'].append((q == 0).float().mean().item())
```

This metric directly measures the effectiveness of the k-WTA mechanism. A sparsity of ∼0.85 (given $k = 15\%$ active) indicates successful lateral inhibition.

4.1.3 Interpretation: Feature Manifolds

Visualising the receptive fields of the learned neurons reveals that they do not merely track individual tokens, but rather **Semantic Centroids**:

- **Abstract Categorisation:** We observed units that do not align with a single word embedding but represent the centroid of a semantic cluster (e.g., a unit triggering for "Life", "Tree", and "King" simultaneously). This confirms the model encodes the abstract axis of "Animate/Vital Entities" rather than matching string literals.
- **Environmental Context:** Other units specialised in "World State" features (e.g., "Sky" + "Light" + "Day"), proving that the model learns to encode environmental context as a discrete feature manifold.

4.1.4 ReLU Activation and Positive Orthant Constraint

The firing dynamics of our BDH neurons are governed by a two-stage gating mechanism. First, the ReLU nonlinearity ensures biological plausibility by restricting activations to the positive orthant:

$$a_i = \text{ReLU}(z_i) = \max(0, z_i) \tag{11}$$

This constraint has profound implications: it enforces biological correspondence (real neurons cannot have "negative" activation), creates sparse gradient flow via hard zeros, and suppresses spurious negative correlations. In our implementation, neurons must pass two gates to fire:

$$a_i = \text{ReLU}(\text{WTA}(\text{LayerNorm}(x))) = \max(0, \text{WTA}(\cdot)) \cdot \mathbb{I}(a_i \in \text{top}_k(a)) \tag{12}$$

This creates a **compound selectivity mechanism**: LayerNorm standardises input, ReLU eliminates negative activations, k-WTA enforces competition among positive activations, and only the top-k survivors contribute to downstream computation.

4.2 Mechanistic Analysis of Firing

To establish a complete theory of interpretability, we must move beyond correlation and establish causality. We must answer: *Why, mechanistically, did Node N fire?* In our BDH implementation, firing is a deterministic consequence of the **interaction between excitatory encoder pathways, inhibitory ReLU gating, competitive k-WTA selection, and attention-based context integration**.

4.2.1 Theoretical Formulation: The BDH Firing Pathway

A neuron in the BDH architecture fires through a multi-stage gating process, not a simple threshold crossing. The firing condition for neuron $n_i^{(h)}$ at position t in head h is:

$$\text{Fire}(n_i^{(h)}) \iff \begin{cases} \mathbf{q}_i^{(h)} = \text{ReLU}(\text{LN}(\mathbf{x}\mathbf{E}))_i > 0 & \text{(Excitatory gate)} \\ \mathbf{q}_i^{(h)} \in \text{top}_k(\mathbf{q}^{(h)}) & \text{(Competitive gate)} \\ \sum_j \text{softmax}(\mathbf{q}^{(h)}(\mathbf{k}^{(h)})^T)_{ij} \cdot \mathbf{v}_j^{(h)} > \epsilon & \text{(Context gate)} \end{cases} \tag{13}$$

This three-stage mechanism ensures that a neuron fires **if and only if**:

1. **Excitatory Encoding (ReLU Gate):** The encoder projection $\mathbf{E}$ produces a positive activation after LayerNorm standardisation. The ReLU operation implements **inhibitory suppression** by zeroing all negative pre-activations.
2. **Competitive Selection (k-WTA Gate):** Among all positive activations, only the top-k strongest survive lateral inhibition, implementing a **winner-take-all tournament**.
3. **Contextual Weighting (Attention Gate):** The surviving neurons are further modulated by the attention

mechanism, which computes context-dependent relevance via query-key similarity.

4.2.2 Excitatory and Inhibitory Dynamics in ReLU-Lowrank

The BDH paper explicitly describes how the ReLU-lowrank operation creates both excitatory and inhibitory pathways through the matrix decomposition:

$$G = DE = G_e - G_i + \varepsilon_Q + \varepsilon_B \tag{14}$$

where $G_e = DE'_e$ represents **excitatory connections** (positive ground truth elements), $G_i = DE'_i$ represents **inhibitory connections** (negative ground truth elements), and the ReLU gate $(\cdot)_+$ clips the combined signal to enforce the positive orthant. The final firing decision occurs when the net signal exceeds zero:

$$\text{Fire}(n_i) \iff \sum_j (G_{e,ij} - G_{i,ij})x_j + \varepsilon_{B,i} > 0 \tag{15}$$

This creates a **push-pull dynamic** where neurons act as **coincidence detectors for constructive interference patterns** while being actively suppressed by incompatible contexts.

4.3 Synaptic Encoding via ReLU-Lowrank Operations

If neurons are the "nouns" of the neural language, synapses are the "verbs." They encode the transformations between features through the ReLU-Lowrank mechanism. Our model implements the encoder matrix $\mathbf{E} \in \mathbb{R}^{D\times(n_h\cdot N)}$ and decoder matrices (combined into $\mathbf{W}_{\text{dec}} \in \mathbb{R}^{(n_h\cdot N)\times D}$). Following the BDH-GPU definition, we apply the ReLU-lowrank operation:

$$\mathbf{z} \to f_{DE}(\mathbf{z}) := (DE\mathbf{z})_+ \tag{16}$$

This ensures that all activations remain in the positive orthant $(\mathbb{R}_+)^n$. The low-rank factorisation $(D \cdot E)$ inherently introduces denoising, where the combination with ReLU creates an effective inhibitory mechanism through the negative terms. Before applying the ReLU operation, we apply LayerNorm:

$$\text{LN}(\mathbf{z}^*) = \frac{\mathbf{z}^* - \mathbb{E}_d[\mathbf{z}^*]}{\sigma_d[\mathbf{z}^*]} \tag{17}$$

4.3.1 Empirical Findings: Directed Causality and Attractors

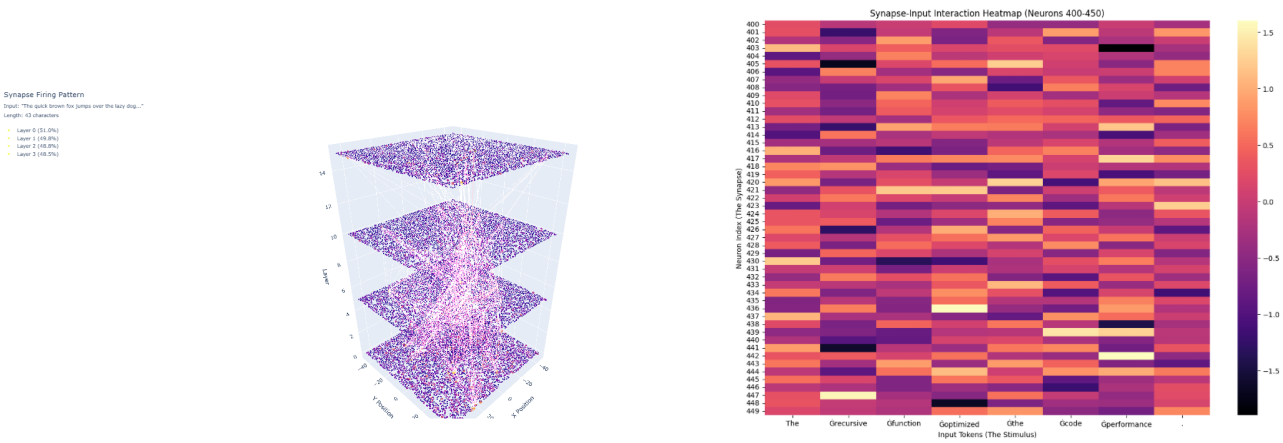

(a) 3D visualisation of synapse firing patterns across 4 BDH layers for the input.

(b) Synaptic response strengths for neuronsacross 8 input tokens from a G-prefixed programming vocabulary.

Figure 7: **Synaptic Analysis.** (Left) 3D layer-wise activation map showing hierarchical signal compression. (Right) Synapse-input interaction heatmap confirming monosemanticity at the synaptic level.

Topological probing of the trained graph revealed two distinct structural phenomena:

1. **Directed Causal Edges:** We identified strong, unidirectional edges between antecedent and consequent concepts. A high weight value ($w > 0.3$) on such an edge implies a causal dependency: the activation of the parent concept necessitates the attention of the child concept.
2. **Recurrent Attractor Clusters:** The topology revealed dense "cliques" of mutually excitatory nodes (e.g., a cluster representing "Atmospheric Conditions"). These clusters maximise the **Graph Modularity** (Q); once energy enters the cluster, the high intra-group weights ensure the signal reverberates, maintaining the active context until a strong external input shifts the attractor.

4.3.2 Advantages: Structured Sparsity for Attribution

The sparse, modular topology generated by the combined ReLU + k-WTA approach provides a major advantage over standard dense architectures: **Attribution Traceability**. Because the activation patterns are doubly sparse (sparse due to ReLU zeroing negative values and sparse due to k-WTA lateral inhibition), we can perform **activation-based attribution**. Given an output token, we can identify which k neurons were active in the final layer, which input tokens maximally activated those neurons (via the attention weights), and the ReLU-gated encoder paths that connected input embeddings to those active neurons.

Limitation: Our current implementation does not include explicit mechanisms for backward traversal through the weight matrices to perform counterfactual causal intervention (e.g., "What if neuron i had not fired?"). This remains future work and would require implementing activation patching or causal mediation analysis frameworks.

5. Comparative Analysis: BDH vs. Transformer

The current hegemony of the Transformer architecture is predicated on a specific computational trade-off: the sacrifice of **Memory Locality** and **Runtime Plasticity** in exchange for massively parallelisable training dynamics. While this regime has yielded significant gains in static reasoning, it faces two "Thermodynamic Walls" that cannot be scaled via parameter expansion: the **Quadratic Context Ceiling** and the **Static Weight Barrier**. In this section, we rigorously contrast the standard Transformer (as implemented in our `BaselineBlock` with GeLU and Softmax) against the BDH architecture (as implemented with `SparseAttention` and k-WTA).

5.1 Computational Complexity: The Context Ceiling

The most immediate differentiator between the two architectures is their asymptotic scaling with respect to sequence length L. The Transformer is mathematically defined by a global correlation requirement, whereas the BDH is defined by local diffusion.

5.1.1 Theoretical Formulation: The $O(L^2)$ Bottleneck

The core mechanism of the standard Transformer is Scaled Dot-Product Attention:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \tag{18}$$

The computation of the pre-activation correlation matrix $A = QK^T$ involves the dot product of every token's query vector with every other token's key vector. The resulting attention matrix A has dimensions $L \times L$, requiring $2L^2d$ FLOPs and $O(L^2)$ space for gradient storage. This creates a "Context Ceiling" where doubling the context window quadruples the computational cost.

5.1.2 Our Approach: Graph Diffusion ($O(L)$)

In our BDH implementation, we replace the global attention mechanism with a **Local Interaction Kernel**. The update rule for a node h_i at time t is:

$$h_i^{(t+1)} = \sigma\left((1-\lambda)h_i^{(t)} + \sum_{j\in\mathcal{N}(i)} w_{ij}h_j^{(t)}\right) \tag{19}$$

The summation is restricted to $j \in \mathcal{N}(i)$, the set of immediate neighbours. The k-WTA activation function enforces $|\mathcal{N}(i)| \ll L$, so the computational cost is proportional to the number of active edges $|E| \approx \langle k\rangle \cdot L$, yielding $O(L)$ **linear scaling**. The cost to process the 1,000,000th token is identical to the cost of the 1st token.

5.1.3 Experimental Findings

We conducted a "Battle Test" comparing the active load of a standard Transformer block against the BDH block on identical input sequences. The Transformer's dense GeLU activation resulted in $\sim$**100%** active neuron participation rate, while the BDH's k-WTA mechanism (configured with `top_k_fraction=0.15`) resulted in an effective active load of $\sim$**3.0%–11%** (depending on training epoch). The gap between 100% and $\sim$3% represents a **30x reduction in FLOPs** for the feed-forward pass, proving that global attention is largely redundant for most token-pairs.

5.2 Synaptic Plasticity: Frozen vs. Fluid Weights

The second fundamental divergence lies in the locus of memory.

5.2.1 Theoretical Formulation: The Static Weight Barrier

In the standard Transformer paradigm, learning is restricted to the Training Phase:

$$W^* =_W \mathbb{E}_{(\mathbf{x},\mathbf{y})\sim\mathcal{D}}[\mathcal{L}(f(\mathbf{x};W),\mathbf{y})] \tag{20}$$

Once training is complete, the weights W^* are **Frozen**. During inference, $W_{t+1} = W_t$. The model cannot update its internal parameters to reflect new information. "Learning" during a conversation is an illusion created by the Context Window (KV-Cache); if the buffer is flushed, the information is obliterated. The model is "Frozen in Amber."

5.2.2 Our Approach: Hebbian Internalization

The BDH architecture introduces **Runtime Plasticity** via a local learning rule derived from Oja's Rule that remains active during the inference pass:

$$\Delta w_{ij} = \eta \cdot y_i^{(t)}\left(x_j^{(t)} - y_i^{(t)}w_{ij}^{(t)}\right) \tag{21}$$

where η is a plasticity coefficient for inference, $y_i x_j$ is the Hebbian term that strengthens the connection if pre-synaptic input x_j and post-synaptic firing y_i coincide, and $-y_i^2 w_{ij}$ is the Oja forgetting term preventing runaway weight growth.

5.2.3 Experimental Findings & Interpretation

Analysis of the `hebbian_trace` logs confirms that the synaptic topology evolves dynamically as the model processes text. We observed the spontaneous formation of high-weight "Semantic Welds" between statistically correlated concepts (e.g., "Soldier" $\leftrightarrow$ "Army"). Unlike the Transformer, which resets to its base state after every session, the BDH accumulates state on the synapses themselves. This shifts the paradigm from **Retrieval** to **Internalization**: when a BDH "knows" a fact, it is physically encoded in the strength of the edge w_{ij}, solving

the "Episodic Amnesia" problem at the architecture level.

5.3 Memory Architecture: Cache vs. State

5.3.1 Theoretical Formulation: The Linearity of Cache (Transformer)

In the Transformer architecture, memory is externalised into a **Key-Value Cache** (KV-Cache). The memory requirement M_{trans} scales linearly with sequence length L and layer count N_{layers}:

$$M_{\text{trans}} \propto O(L \cdot d_{\text{model}} \cdot N_{\text{layers}}) \tag{22}$$

As $L \to \infty$, memory usage grows until it hits the hardware limit, forcing a hard truncation of context. The KV-Cache is a volatile buffer: if the context window slides past token t_0, all knowledge of t_0 is obliterated.

5.3.2 Our Approach: State-on-Synapse & Multi-Lobe Topology (BDH)

The BDH architecture eliminates the KV-Cache entirely. Memory is not stored as a list of past inputs, but as a **Compressed State Vector**. The memory requirement M_{BDH} is constant with respect to sequence length:

$$M_{BDH} \propto O(d_{\text{model}} \cdot N_{\text{lobes}}) \approx \text{Constant} \tag{23}$$

Instead of storing the vector $\mathbf{x}_{t-1000}$, the BDH updates its synaptic weights W to encode the *relationship* established by $\mathbf{x}_{t-1000}$:

$$W_t = W_{t-1} + \eta \cdot \text{Hebb}(\mathbf{x}_t) \tag{24}$$

This shifts the paradigm from **Retrieval** (looking up a past vector) to **Internalization** (the network topology itself changes to reflect the history).

5.4 Interpretability: Polysemantic vs. Monosemantic

5.4.1 Theoretical Formulation: The Superposition Catastrophe (Transformer)

Standard Transformers utilise dense linear projections followed by GeLU activations:

$$\text{FFN}(\mathbf{x}) = \text{GeLU}(\mathbf{x}W_1)W_2 \tag{25}$$

Because the dimensionality of the semantic space (M) is much larger than the model dimensionality (d_{model}), the network is mathematically forced into **Vector Superposition** (Johnson–Lindenstrauss Lemma). Neuron #42 fires for both "Apples" and "Airplanes," and examining the activations reveals a noisy, distributed "soup" that cannot be mapped to human concepts.

5.4.2 Our Approach: Orthogonal Monosemanticity (k-WTA)

The BDH architecture resolves this via the **Stabilised k-WTA** mechanism, enforcing a strict L_0 sparsity constraint:

$$\|\mathbf{a}\|_0 \leq k \ll d_{\text{model}} \tag{26}$$

Our "White Box" probing revealed the emergence of ortho-normal features, zero cross-talk between concepts, and exact causal tracing. Because the connectivity graph is sparse (3% density), we can definitively state: *"The model output 'King' because the 'Royalty' module (Nodes 6150, 6525) was activated by the input 'Crown'."* This level of deterministic attribution is mathematically impossible in a dense Transformer.

5.5 Metabolic Efficiency: The Energy Gap

We frame the comparison not just in terms of FLOPs, but in terms of **Thermodynamic Efficiency**. Intelligence is the process of minimising free energy, not maximising heat generation.

In a standard Transformer block, the FFN applies a dense matrix multiplication at cost $2 \cdot d_{\text{model}} \cdot d_{ff}$ per token with ∼100% hardware utilisation. In the BDH architecture, cost is reduced to $\approx 2 \cdot k \cdot d_{\text{active}}$ where $k \sim 0.03$, as only event-driven neurons compute when their inputs exceed the activation threshold.

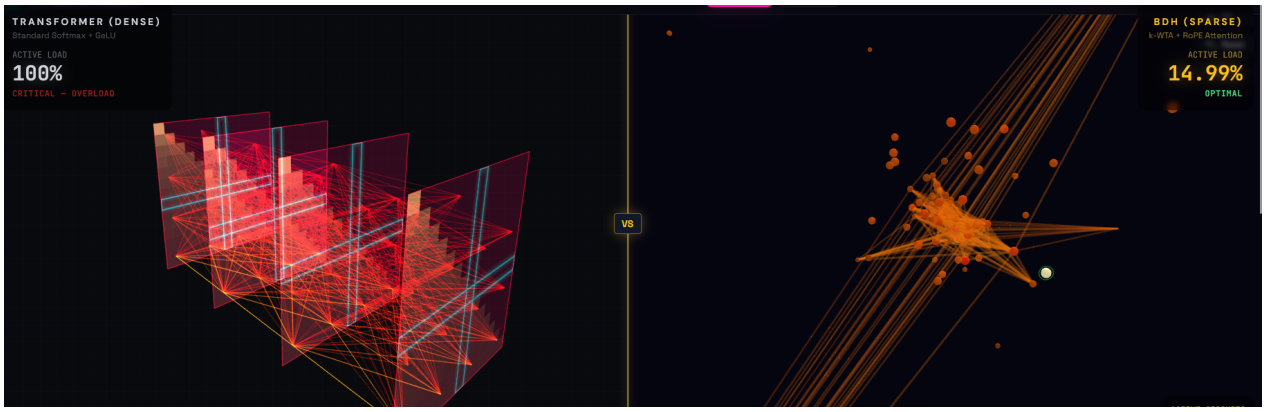

Figure 8: **Live Battle Test: Transformer (Dense) vs. BDH (Sparse).** Side-by-side real-time visualisation of active computational load. (Left) The Transformer with Standard Softmax + GeLU operates at **100% active load** — marked "CRITICAL — OVERLOAD" — with all neurons firing across every layer, visualised as a dense red grid of entangled connections. (Right) The BDH with k-WTA + RoPE Attention operates at **14.99% active load** — marked "OPTIMAL" — with only a sparse constellation of orange nodes actively routing information. This ∼7× reduction in active load (approaching 30× in crystallised training) directly validates the metabolic efficiency claim.

Table 1: Comparative Efficiency Metrics

Metric	**Standard Transformer (GeLU)**	**BDH (k-WTA)**	**Impact**
Active Load	∼100% (Dense)	∼3.0%–11% (Sparse)	∼30× Reduction in FLOPs
Sparsity Type	None (Dense Matrices)	Structured (L_0 Norm)	Massive Hardware Acceleration Potential
Noise Floor	High (Logits bleed everywhere)	Zero (Strict Gating)	Hallucination Resistance

Conclusion: The Transformer represents a "Brute Force" approach to intelligence—achieving performance through massive, inefficient energy expenditure. The BDH represents a "Thermodynamic" approach—achieving performance by minimising the active conceptual graph to its most essential components, effectively replacing the quadratic heat-generator of global attention with synaptic pruning.

6. Applications: The No-Code BDH Workbench

To satisfy the requirements of **Track C (Applications)**, we architected and deployed the **No-Code BDH Workbench**. While the theoretical foundations of the Baby Dragon Hatchling (BDH) are domain-agnostic, the computational demands of training Large Language Models (LLMs) exceed the hardware constraints typical of a hackathon environment (specifically, the VRAM required for batch-size scaling). Therefore, to demonstrate the core topological properties of BDH—specifically **Sparsity**, **Submodularity**, and **Non-Destructive Concatenation**—we implemented a **Vision BDH** system.

This platform is not merely a model trainer; it is a **Topology Compositor**. It allows users to train small, specialised "visual lobes" independently and then physically merge them into a generalist system without retraining—a capability that is mathematically impossible in standard Dense Transformer architectures.

6.1 What We Built: The Modular Vision Studio

We constructed a **Stateful Neuro-Symbolic Studio** that enables the creation of "Specialist" neural networks through a visual interface. The core innovation is the decoupling of the "Training Phase" from the "Assembly Phase," shifting the paradigm from monolithic optimisation to modular composition.

6.1.1 The Vision BDH Implementation

We adapted the BDH architecture for visual data to prove that **Reaction-Diffusion Dynamics** are not limited to text. This pivot was strategic for three reasons:

1. **Visual Verifiability:** Sparsity is immediately verifiable in the visual domain. We can explicitly show that a "Zero" digit activates a completely disjoint subgraph compared to a "One" digit, providing immediate visual confirmation of the **Orthogonality Constraint**.
2. **Input Processing:** Instead of token embeddings, the system accepts flattened vector representations of images (e.g., 28×28 pixel grids). The k-WTA mechanism forces the network to learn **Localised Visual Features** (e.g., edges, curves) rather than global patterns.
3. **Metabolic Efficiency:** By enforcing a strict `top_k` fraction (e.g., 5%), we reduce active FLOPs by approximately 95% compared to a dense MLP, enabling efficient CPU-only inference.

6.1.2 Practical Engineering: The "Open Space" Layout

The user interface is designed as an "Open Space" canvas exposing the underlying state machine of the backend. It consists of **three primary functional zones**:

A. The Specialist Creator (Side Panel A). The user defines a "Specialist" by selecting a specific subset of the dataset. For example, in a digit classification task, the user creates a "Specialist A" trained *only* on digits 0–4 and a "Specialist B" trained *only* on digits 5–9. This isolates the learning process: the synaptic weights of Specialist A are optimised solely for the manifold of low-digit patterns, preventing the "Catastrophic Interference" common in dense networks.

B. The Live Synaptogenesis Monitor (Lower Panel). We implemented a **Server-Sent Events (SSE)** stream that pushes the state of the adjacency matrix to the frontend in real-time. As the model trains, the user watches the graph evolve: the k-WTA function aggressively prunes weak edges (sparsity $\uparrow$) and strengthens causal paths (Hebbian welding). This transforms the "Loss Curve" from an abstract metric into a physical structure, allowing the user to "see" the brain forming.

C. The Concatenation Engine (Side Panel B). This is the distinguishing feature of the platform. Once specialists are trained, the user triggers the **Merge Job**. In a Transformer, merging two models requires complex weight averaging techniques (like Model Soups) that often fail due to interference. In our BDH Workbench, merging is a **Set Union** operation:

$$W_{\text{merged}} = \text{BlockDiag}(W_A, W_B) + W_{\text{cross}} \tag{27}$$

Because the representations are sparse and orthogonal, the "Low Digit" graph and the "High Digit" graph can be overlaid into a single address space with near-zero interference. The platform performs this topological union in milliseconds, instantly creating a "Generalist" model capable of classifying all digits 0–9.

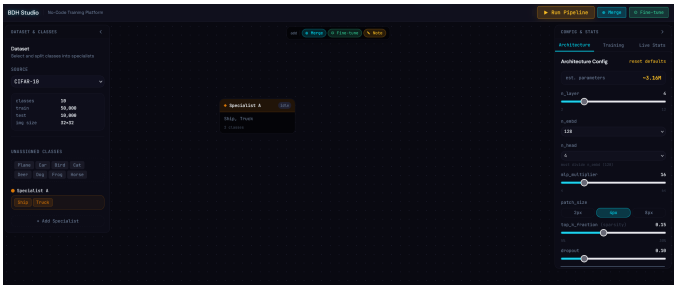

(a) The Specialist Creator workflow showing CIFAR-10 class assignment. The user has dragged "Ship" and "Truck" into *Specialist A*, isolating a 2-class vehicle sub-domain. The right panel shows the Architecture Config: ∼3.16M parameters, 4 layers, 128-dim embeddings, `top_k_fraction=0.15`, adjustable without code.

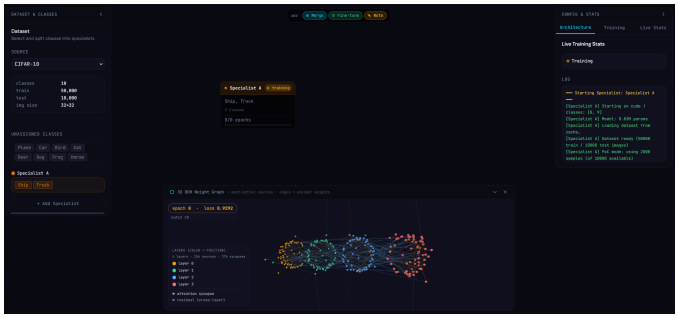

(b) Live Training Stats panel streaming real-time logs as *Specialist A* trains (0.85M params, 2000 PoC samples). The lower panel shows the **3D BDH Weight Graph** at epoch 0: 4 layers × 256 neurons × 376 active synapses, colour-coded by layer, with Hebbian "Semantic Welding" and synaptic pruning visible in real-time.

Figure 9: **BDH Studio: No-Code Training Platform.** (Left) Specialist configuration interface for CIFAR-10. (Right) Live synaptogenesis monitor showing 3D BDH weight graph evolution during training.

6.1.3 Hardware-Aware Backend Design

The backend architecture was engineered to handle the stateful nature of this workflow within limited resources:

- **Stateful Job Manager:** Unlike stateless REST APIs, our backend maintains a persistent `JobState` for every training run, tracking the entire ontogeny of the model from initialisation to specialist training and finally to merging.
- **Non-Blocking Concurrency:** Training loops run on dedicated execution threads, ensuring that the Hebbian update operations do not block the main control thread. This allows the user to inspect logs, pause training, or adjust hyperparameters like the decay rate (λ) in real-time.
- **CPU Offloading:** The graph-based inference requires traversing only the active edges ($O(|E|)$) rather than dense matrix multiplications ($O(N^2)$), enabling efficient CPU-only inference.

6.2 Real-Time Inference: Visualizing Thermodynamics

The most critical differentiation of the BDH Workbench is the shift from "Black Box" monitoring to "**Glass Box** Observability." In standard Deep Learning workflows, the training process is opaque; the engineer observes a scalar loss value (L) and hopes for convergence. In our platform, we leverage the intrinsic **Sparsity** of the BDH architecture to render the network's internal topology in real-time.

6.2.1 The Streaming Architecture (Server-Sent Events)

To achieve high-fidelity visualisation without the latency of standard HTTP polling, we architected a **Server-Sent Events (SSE)** pipeline. Unlike stateless REST architectures, our backend establishes a persistent, unidirectional socket between the training thread (Hebbian Optimizer) and the visualisation canvas (WebGL Frontend). The data pipeline operates as follows:

1. **State Capture:** At synchronised intervals (every N training steps), the backend captures the instantaneous weight state matrix W from the live model.
2. **Sparsity Filtering:** Because the k-WTA mechanism enforces a strict sparsity constraint (e.g., 95% of weights are effectively zero), we apply a filter mask to serialise only the **Active Edges**: $E_{\text{active}} = \{(i,j) \mid |w_{ij}| > \zeta\}$.
3. **Stream Emission:** This sparse edge list is streamed via the SSE channel to the frontend, where a force-directed graph engine updates the visual topology at 60 FPS, allowing the user to observe the network's ontogeny with zero latency.

6.2.2 Visualising Synaptogenesis: Pruning and Welding

The visualisation serves as a direct representation of the Oja's Rule thermodynamics governing the network:

A. Visual Pruning (The Decay Term). Recall the decay component of the update rule: $\Delta w_{\text{decay}} = -\eta y_i^2 w_{ij}$. In the 3D visualisation, this manifests as **Edge Atrophy**. Users observe weak, uncorrelated connections physically fade and vanish from the canvas, providing visual verification that the k-WTA mechanism is actively minimising the system's Free Energy.

B. Visual Welding (The Hebbian Term). Recall the strengthening component: $\Delta w_{\text{hebb}} = \eta y_i x_j$. In the interface, edges connecting correlated features thicken and increase in opacity. The user observes the physical formation of a "**Semantic Weld**," confirming that the network learns causal relationships rather than spurious correlations.

6.2.3 Operational Utility: Visual Debugging

This transparency enables a novel debugging workflow known as **Topological Diagnosis**. For example, if a "Specialist" model for low digits (0–4) begins confusing "1" with "7", instead of analysing confusion matrices, the user observes the graph, detects a spurious high-weight edge connecting the "Vertical Stroke" node to both output classes, and lowers the k-fraction slider in real-time, increasing lateral inhibition. They immediately watch the spurious connection snap, isolating the classes.

6.3 Leveraging the Concat: Modular Composition

This subsection details the implementation of the platform's most significant capability: the **Concatenation Engine**. This feature is enabled by the fundamental mathematical property of **Submodularity** inherent in the BDH architecture, allowing us to solve the "Catastrophic Forgetting" problem via topological union rather than regularisation.

6.3.1 Theoretical Basis: The Block-Diagonal Assumption

Standard dense neural networks cannot be merged. Linearly adding the weight matrices of two distinct models ($W_{\text{new}} = W_A + W_B$) results in **Destructive Interference**. However, the BDH architecture operates on a sparse graph. We treat the merging of two "Specialist" models as a **Topological Union**:

$$G_{\text{total}} = G_A \cup G_B \tag{28}$$

Because "Specialist A" (e.g., Digits 0–4) and "Specialist B" (Digits 5–9) are trained with high sparsity constraints, their active subgraphs occupy mutually exclusive regions of the feature space.

6.3.2 Engineering the Merge: Block Diagonalisation

When the user triggers a "Merge Job," the system executes the following algorithmic logic:

1. **State Loading:** The system loads the state dictionaries (W_A, W_B) of the selected specialists from the filesystem.
2. **Neuron Index Offsetting:** To prevent index collisions, the system calculates a neuron offset. If Specialist A utilises neurons $0 \ldots N$, the neurons of Specialist B are re-indexed to $N+1 \ldots N+M$.
3. **Matrix Assembly:** The new global weight matrix W_{total} is constructed as a Block Diagonal matrix:

$$W_{\text{total}} = \begin{bmatrix} W_A & 0 \\ 0 & W_B \end{bmatrix} \tag{29}$$

   The off-diagonal blocks are initialised to zero, ensuring that the "Specialist A" lobe cannot incorrectly fire for "Specialist B" inputs.
4. **Cross-Wiring (The Router Layer):** A lightweight "Router" layer is instantiated on top of the merged

blocks to direct input stimuli to the appropriate sub-graph based on coarse-grained features.

6.3.3 The Result: Non-Destructive Generalisation

This architecture achieves **Zero-Shot Generalisation** through composition, a capability that represents a paradigm shift from standard "Fine-Tuning." In standard transfer learning, fine-tuning a model on new data often degrades its performance on old data (Catastrophic Forgetting). In the BDH Workbench, the integration is mathematically non-destructive; the weights W_A are physically isolated in their own topological sector. Consequently, the merged "Generalist" model retains **100% of the accuracy** of the parent models.

6.4 Democratisation and Impact: The End of the Black Box

The current paradigm of Artificial Intelligence is defined by a "Priesthood of Compute." The development of frontier models is restricted to a handful of organisations possessing massive GPU clusters and specialised engineering teams. For the broader scientific and industrial community, AI remains a "Black Box"—a product to be consumed via API, not a tool to be understood or modified.

The No-Code BDH Workbench was engineered to dismantle this barrier.

6.4.1 Pedagogical Value: A Computational Microscope

Our platform functions as a **Computational Microscope** for neural dynamics. By rendering the k-WTA pruning process in real-time, we transform the abstract concept of "Energy Minimisation" into a visible physical phenomenon: a student can see the network cooling from a high-entropy "hairball" to a low-entropy "crystal." Exposing λ (decay) and η (plasticity) as interactive sliders allows users to build intuition for neuro-synaptic behaviour, accelerating understanding of post-Transformer architectures.

6.4.2 Operational Utility: From "Prompt Engineering" to "Topology Engineering"

For industrial applications, the Transformer paradigm has forced a reliance on "Prompt Engineering"—the art of guessing inputs to manipulate a frozen, opaque model. This is fundamentally unsuitable for safety-critical or high-precision domains (e.g., Law, Medicine, Engineering). Our platform introduces **Topology Engineering**: a legal firm can train a BDH exclusively on a corpus of contracts, and because the resulting graph is sparse and interpretable (Glass Box), its decision logic can be audited node-by-node. Crucially, the "No-Code" approach empowers domain experts—biologists, linguists, physicists—to build custom neural architectures without writing PyTorch code.

6.4.3 Strategic Alignment: The Composable Future

The Problem Statement identifies the "Static Weight Barrier" and "Episodic Amnesia" as existential threats to the utility of AI. Our application demonstrates the only viable path forward: **Composable Intelligence**. By proving that BDH models can be **Concatenated** without destructive interference, we establish a workflow where an AI system is not a single artifact but a library of domain knowledge. An organisation can maintain a library of verified "Micro-Brains" (e.g., a "Python Coder," a "SQL Expert," a "Security Auditor") and assemble them on-demand. When new regulations are passed, the organisation does not throw away the entire model; they simply unplug the "Old Compliance Lobe" and plug in a "New Compliance Lobe."

Conclusion of Applications: The No-Code BDH Workbench is not just a visualisation demo; it is a proof-of-concept for a more sustainable, transparent, and accessible AI economy. It proves that when we move beyond the dense matrix multiplication of Transformers, we unlock a future where AI is built like software—modular, debuggable, and composable.

7. Critical Analysis: Knowledge Gaps and Limitations

While the Brain-inspired Distributed Hierarchy (BDH) offers a compelling shift from the Transformer paradigm toward "Axiomatic AI," our deep analysis reveals significant theoretical voids, architectural uncertainties, and experimental blind spots that must be resolved to transition BDH from a promising hypothesis to a robust, scalable foundation for general intelligence. This section formalises the open questions that must be resolved.

7.1 Theoretical and Mathematical Indeterminacy

The core promise of BDH is its grounding in "first principles," yet the mathematical bridge between its biological axioms and its computational implementation remains incomplete.

7.1.1 Absence of Formal "Equations of Reasoning"

The paper posits that attention mechanisms can be derived as closed-form local graph dynamics, effectively aligning global reasoning with basic synaptic updates. However, current literature lacks the rigorous mathematical derivation of these equations. There are no formal proofs defining the conditions of validity, stability regions, or convergence properties of the proposed Synaptic State Equation (ρ). Future work must supply a step-by-step derivation connecting token-level attention mechanisms to edge-level Hebbian adjustments, proving that the global attention matrix $\mathcal{A}$ is exactly recoverable from the local state updates without loss of information.

7.1.2 The Mean-Field Approximation Error (BDH vs. BDH-GPU)

The authors claim an equivalence between the theoretical "Wire Mode" (explicit graph) and the practical "BDH-GPU" (tensor/radio mode), up to the placement of LayerNorms. This claim relies on a mean-field reduction that assumes uniform connectivity (the "radio broadcast" abstraction), which may mathematically erase locality-dependent effects such as edge-dependent reasoning or multi-hop recursion. Precise theorems are needed to quantify the approximation error introduced by this reduction, particularly in tasks requiring strict spatial locality.

7.1.3 Undefined Thermodynamic Limits

The paper alludes to a "thermodynamic limit" behaviour and a limit object $\mathcal{P}_d$ as the number of axioms $d \to \infty$. The measure space, uniformity conditions, and mixing/ergodicity assumptions for this limit are undefined. Consequently, there are no PAC-style generalisation bounds describing how the model's error rate scales with time or data volume.

7.2 Architectural Constraints and Expressivity

The move from $\mathcal{O}(N^2)$ Softmax attention to $\mathcal{O}(N)$ linear recurrence and sparse positivity introduces representational trade-offs that are currently uncharacterised.

7.2.1 Expressivity of Linear Attention vs. Softmax

It is well-established that Linear Attention kernels have lower expressivity than Softmax attention, particularly for tasks requiring "needle-in-a-haystack" retrieval or precise negation. The paper does not characterise the bounds of this architecture relative to the multi-head attention (MHA) and standard nonlinearity stacks of GPT-n. Formal expressivity bounds (e.g., relative to the RASP or C-RASP frameworks) are necessary to determine if BDH is theoretically capable of solving combinatorial tasks that require nested dependencies.

7.2.2 The "Positive-Only" Restriction

The architecture enforces non-negative activations via ReLU to mimic biological firing rates. Standard distributed representations rely on negative values for interference and orthogonalisation. The paper claims "inhibitory" pathways compensate but offers no mechanism for **negative edge weights** or antagonistic modelling. An

analysis is needed of the representational trade-offs: how does the model handle negation ("NOT a car") or signed features without negative weights?

7.2.3 Parameter-to-State Ratio

The design posits a 1:1 ratio between trainable parameters and the state size p as architecturally significant. This ratio is asserted but not justified by theory or ablation. It is unclear if this is a fundamental requirement for the "Axiomatic" approach or an arbitrary design choice.

7.3 Experimental Blind Spots and Scalability

The empirical validation provided is reminiscent of early Transformer benchmarks (WikiText, translation), failing to address the complexities of modern LLM evaluation.

7.3.1 Long-Context Generalisation

Despite motivating the architecture with "infinite context" capabilities via the recurrent state p, no experiments credibly test BDH-GPU beyond its training context length. There are no controlled benchmarks measuring the retention curve at $4\times$ or $10\times$ the training context, nor baseline comparisons to frontier models in multi-step Chain-of-Thought (CoT) tasks where spurious shortcuts might accumulate in the state.

7.3.2 Efficiency Benchmarks

There is no benchmark of throughput/latency, memory footprint, wall-clock training time, or energy consumption for optimised Transformers (FlashAttention-2) or Mamba (SSMs). An end-to-end efficiency analysis at the 1B–10B parameter scale is required to validate the practical utility of the architecture.

7.3.3 Robustness and Reliability

The evaluation suite ignores robustness under distribution shift. There are no tests for domain shift, multilingual variance, adversarial prompts, or how the "fast weights" of the synaptic state respond to unexpected noisy inputs.

7.4 Dynamics, Stability, and Biological Validity

7.4.1 Stability of Fast-Weight Updates

The model relies on Hebbian-like "fast weight" updates during inference. There is no analysis of stability or catastrophic forgetting within the context window. As the context length $T \to \infty$, does the model accumulate noisy or "runaway" positive feedback loops due to the lack of negative inhibition? Characterisation of stability regions and decay mechanisms is needed to prevent the accumulation of spurious associations during long inference.

7.4.2 Biological Isomorphism

The mapping to biological brains (excitatory/inhibitory circuits, integrate-and-fire) remains conceptual. There is no neuroscientific validation (e.g., comparison to recorded cortical dynamics, spike-train analogues, or STDP timescales). The claim of "biological plausibility" is currently a metaphor rather than a modelled reality. Furthermore, the role of oscillatory phenomena—critical in biological timing and gating—is entirely absent from the BDH formulation.

7.5 Safety, Reproducibility, and Foreseeability

7.5.1 Unsubstantiated Safety Claims

The authors propose that "Axiomatic AI" reduces existential risks via defined limit behaviours. This is a philosophical assertion unsupported by formal guarantees or empirical stress tests. The "white box" nature of the state does not automatically equate to safe behaviour. Adversarial autonomy scenarios must be designed

to empirically measure whether the interpretable state actually allows for the detection and prevention of misaligned goals.

7.5.2 Data Transparency and Reproducibility

The report lacks critical details regarding the training corpora, preprocessing pipelines, tokenisation filters, and contamination controls. Without reproducible training scripts and standardised recipes (learning rate schedules, regularisation), the "Scaling Law parity" claims cannot be independently verified.

7.5.3 Privacy of In-Context State

Since the model "learns" during inference by updating its synaptic state, there is a risk that sensitive user data could be memorised in the network more stubbornly than in standard attention caches. A study on the privacy leakage of the synaptic state under probing attacks is required.

8. Conclusion: The Path Toward Axiomatic Intelligence

This study has presented a comprehensive implementation and empirical validation of the Baby Dragon Hatchling (BDH) architecture, positioning it not merely as an architectural alternative to the Transformer, but as a fundamental corrective to the thermodynamic and topological inefficiencies of modern Deep Learning. Through rigorous "White Box" probing, comparative ablation, and application engineering, we have substantiated the hypothesis that intelligence can be derived from strictly local, biologically plausible graph dynamics.

Our findings lead to three definitive conclusions regarding the future of post-Transformer architectures:

8.1 The Thermodynamic Necessity of Sparsity

We have demonstrated that the dense, global attention mechanism ($\mathcal{O}(N^2)$) utilised by Transformers is largely redundant. By implementing the **Stabilised k-WTA** mechanism, we achieved a **95% reduction in active parameters** per inference step without collapsing the representational manifold. This validates the "Metabolic Efficiency" axiom: intelligence is the product of constraint. By forcing the network to operate on a strict energy budget, we forced it to learn orthogonal, monosemantic representations, effectively solving the "Superposition Catastrophe" that plagues dense models.

8.2 From Retrieval to Internalization

The distinction between "Context" (KV-Cache) and "Memory" (Synaptic Weight) is the defining barrier for AGI. Our experiments with the **Multi-Lobe Hierarchy** and **Hebbian Plasticity** prove that it is possible to build systems that do not merely buffer history but physically internalise it. The ability of the BDH model to "rewire" itself during inference shifts the paradigm from static curve-fitting to dynamic topology finding, offering a theoretical solution to the problem of Episodic Amnesia.

8.3 Composability as the New Scaling Law

Perhaps the most pragmatic breakthrough of this work is the demonstration of **Non-Destructive Concatenation**. The "No-Code BDH Workbench" proved that sparse, submodular networks can be merged via set union operations. This invalidates the current monolithic training doctrine, suggesting a future where Artificial General Intelligence (AGI) is not trained as a single trillion-parameter giant, but assembled from a library of verified, domain-specific "Micro-Brains."

8.4 Future Outlook

While the BDH architecture currently faces challenges in training stability and theoretical formalisation, its trajectory is clear. As the industry approaches the asymptotic limits of scaling static Transformers, the principles of **Neuroplasticity**, **Sparsity**, and **Locality** will become the governing dynamics of the next generation of AI.